

L50 - Lab 5, Kernel Tracing 2 & Performance with Antagonists

Professor Andrew W. Moore and Dr Gary Guo

Lent, 2022/2023

The aim of this lab is

- to study application performance using the KUtrace system alongside the application integrated logging used in previous labs, and
- to extend this work by studying the impact of an antagonistic workload upon the previous lab workloads.

The approach for this lab will be to rerun a number of experiments from Lab 3 – the `server_disk` experiments of Book Section 7.7 (Experiments 3 and 4) and then to study this application using the KUtrace system.

1 Tracing `server_disk`

Run the `server_disk` write example under KUtrace

Where required, refer to the lab4 handout for instructions on loading the KUtrace module and setting the size of the trace buffer.

Refer to the lab3 handout for instructions on operating experiments 3 and 4 from that laboratory.

Where necessary load the trace with a large trace buffer, e.g.

```
sudo insmod kustrace_mod.ko tracemb=40
```

Beware the resulting html file will be extremely large.

These traces can be quite long, by running KUtrace without instruction counter information, the trace will be more memory efficient, e.g.

```
~/KUtrace/control/control/kutrace_control
Entering kutrace_control
control> gowrap <<<- this command starts the trace in flight-recorder mode
```

Remember you are studying server behaviour so run KUtrace on the server host.

When running the `server_disk` checksum example, a 20MB of trace may be sufficient.

For each of these experiments, identify regions of particular, peculiar, unusual, or unexpected patterns.

2 Operating with an Antagonist

The use of an antagonist can identify specific problem areas; by putting a component under pressure (memory, CPU, disk, network, etc.) we can identify both immediate issues and points with potential for starvation.

For these experiments repeat Lab3's experiment 3 and 4 again but this time with multiple copies of an antagonist. Our antagonist will be a `cpuhog` in `~/KUtrace/bookcode/aw_files/book-user-code`

Code for the `cpuhog` can be recompiled using the command:

```
(code in ~/KUtrace/bookcode/aw_files/book-user-code/cpuhog_highissue.cc)
g++ cpuhog_highissue.cc -o cpuhog
```

`cpuhog` will exit after about 50 minutes.

You may uncover interesting behaviours as you run experiments with varying numbers of simultaneous `cpuhog` executables.

An example command to run three `cpuhog` processes with the server.

```
~/KUtrace/bookcode/aw_files/book-user-code/cpuhog & \
~/KUtrace/bookcode/aw_files/book-user-code/cpuhog & \
~/KUtrace/bookcode/aw_files/book-user-code/cpuhog & \
~/KUtrace/bookcode/aw_files/book-user-code/server_disk diskdir
```

Further CPU pressure can be achieved by increasing the number of `cpuhog` processes and by changing the scheduling priority. This is an example using the `nice(1)` command to increase run priority for each of the three `cpuhog` processes. You may wish to experiment with other combinations to trigger an interesting (study-worthy) effect.

```
cat diskdir/a* diskdir/c* diskdir/d* > /dev/null
sudo nice --10 ~/KUtrace/bookcode/aw_files/book-user-code/cpuhog & \
sudo nice --10 ~/KUtrace/bookcode/aw_files/book-user-code/cpuhog & \
sudo nice --10 ~/KUtrace/bookcode/aw_files/book-user-code/cpuhog & \
~/KUtrace/bookcode/aw_files/book-user-code/server_disk diskdir
```

For each set of experiments, recall the commands to purge and reload the operating system block cache; this will provide better repeatability and also by starting in a known state, your experiments become more repeatable too.

```
# Purge page cache
echo 1 | sudo tee /proc/sys/vm/drop_caches
# Bring aaaa and dddd back to RAM - suitable for experiment 3
cat diskdir/a* diskdir/d* > /dev/null
```

Check the lab3 handout for further details.

3 Understanding Your Experiments

Remember to start each experiment with a model of your understanding, then run the experiment and you can either validate or improve, develop, and refine your understanding.

By this stage you will have run multiple experiments and followed these observations with exploration of behaviour varying as it is subject to antagonistic behaviour from cpuhog processes. Identify behaviours you have been alerted to in the textbook, three examples can include lock contention, cache misuse, or disk contentions; other issues are also present. These issues and further issues will develop when subjected to cpu contention and cpu exhaustion.

It is precisely such situations that lead to poorer behaviour and that may be identified with tools such as KUtrace. For each experiment look to create (in the antagonistic) or simply to identify one such performance issues, to document the issues and to clearly demonstrate your understanding of the wealth of tracing information presented.

A single lab report will be required for the final two labs. Instructions for the final lab report will be provided shortly.

The following prompts are intended to help you understand your results, while you also provide supporting evidence for your report. However, they are just suggestions — feel free to approach the data differently!

- Discuss the methodology of the measurement tools.
- Explain how the limitations of the methodology are mitigated in this lab.
- Explore the limitations of the experiments conducted in this lab, and explain where the quality of the experiment (e.g., setup, methodology) could have been improved or altered.

Finally, you should always look for odd or surprising results (such as time spent in unexpected ways), and try to explain them. Note that sometimes exceptional results indicate a problem in your setup or scripts.